

Q1. Generate a list of 100 integers containing values between 90 to 130 and store it in the variable `int_list` . After generating the list, find the following: (i) Write a Python function to calculate the mean of a given list of numbers. Create a function to find the median of a list of numbers?

Answer =>

```
import numpy as np
np.random.seed(0)
int_list = np.random.randint(90, 131, 100)
print(int_list)

def mean_list(lst):
    return sum(lst) / len(lst)

print("Mean:", mean_list(int_list))

def median_list(lst):
    lst = sorted(lst)
    n = len(lst)
    mid = n // 2
    if n % 2 == 0:
        return (lst[mid - 1] + lst[mid]) / 2
    else:
        return lst[mid]

print("Median:", median_list(int_list))
```

(ii) Develop a program to compute the mode of a list of integers.

Answer:

```
from collections import Counter
```

```
def mode_list(lst):
    freq = Counter(lst)
    max_count = max(freq.values())
    mode_values = [k for k, v in freq.items() if v == max_count]
```

```
return mode_values

print("Mode:", mode_list(int_list))
```

(iii) Implement a function to calculate the weighted mean of values and weights.

Answer:

```
import numpy as np
```

```
def weighted_mean(values, weights):
    values = np.array(values)
    weights = np.array(weights)
    return np.sum(values * weights) / np.sum(weights)
```

```
weights = np.random.randint(1, 6, 100)
print("Weighted Mean:", weighted_mean(int_list, weights))
```

(iv) Write a Python function to find the geometric mean of a list of positive numbers.

Answer:

```
def geometric_mean(lst):
    product = 1
    for x in lst:
        product *= x
    return product ** (1 / len(lst))
```

```
print("Geometric Mean:", geometric_mean(int_list))
```

(v) Create a program to calculate the harmonic mean of a list of values.

Answer:

```
def harmonic_mean(lst):
    return len(lst) / sum(1/x for x in lst)

print("Harmonic Mean:", harmonic_mean(int_list))
```

(vi) Build a function to determine the midrange of a list of numbers.

Answer:

```
def midrange(lst):
    return (min(lst) + max(lst)) / 2
```

```
print("Midrange:", midrange(int_list))
```

(vii) Implement a Python program to find the trimmed mean of a list, excluding a certain percentage of outliers.

Answer:

```
def trimmed_mean(lst, trim_percent=0.1):
    lst = sorted(lst)
    n = len(lst)
    trim_count = int(n * trim_percent)
    trimmed_lst = lst[trim_count : n - trim_count]
    return sum(trimmed_lst) / len(trimmed_lst)
```

```
print("Trimmed Mean (10%):", trimmed_mean(int_list, 0.1))
```

Q2. Generate int_list2 (500 integers between 200 to 300).

Answer:

```
import numpy as np
np.random.seed(0)
```

```
int_list2 = np.random.randint(200, 301, 500)
print(int_list2)
```

i) Compare the given list of visualization for the given data

1. Frequency & Gaussian distribution

Answer:

```
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats
plt.figure(figsize=(7,5))
sns.histplot(int_list2, bins=25, stat="density")
x = np.linspace(min(int_list2), max(int_list2), 200)
plt.plot(x, stats.norm.pdf(x, np.mean(int_list2), np.std(int_list2)), linewidth=2)
plt.title("Frequency & Gaussian Distribution")
plt.show()
```

2. Frequency smoothed KDE plot

Answer:

```
plt.figure(figsize=(7,5))
sns.histplot(int_list2, bins=25, kde=True, stat="density")
plt.title("Frequency Smoothened KDE Plot")
plt.show()
```

3. Gaussian distribution & smoothed KDE plot

Answer:

```
plt.figure(figsize=(7,5))
sns.kdeplot(int_list2, linewidth=2)
x = np.linspace(min(int_list2), max(int_list2), 200)
plt.plot(x, stats.norm.pdf(x, np.mean(int_list2), np.std(int_list2)), linewidth=2)
```

```
plt.title("Gaussian Distribution & Smoothened KDE Plot")
plt.show()
```

(ii) Write a Python function to calculate the range of a given list of numbers.

Answer:

```
def data_range(lst):
    return max(lst) - min(lst)
print("Range:", data_range(int_list2))
```

(iii) Create a program to find the variance and standard deviation of a list of numbers.

Answer:

```
def variance_std(lst):
    arr = np.array(lst)
    variance = np.var(arr, ddof=1)
    std_dev = np.std(arr, ddof=1)
    return variance, std_dev
```

```
var, std = variance_std(int_list2)
print("Variance:", var)
print("Standard Deviation:", std)
```

(iv) Implement a function to compute the interquartile range (IQR) of a list of values.

Answer:

```
def iqr(lst):
    q1 = np.percentile(lst, 25)
    q3 = np.percentile(lst, 75)
    return q3 - q1
print("IQR:", iqr(int_list2))
```

(v) Build a program to calculate the coefficient of variation for a dataset.

Answer:

```
def coefficient_variation(lst):
    arr = np.array(lst)
    return (np.std(arr, ddof=1) / np.mean(arr)) * 100
print("Coefficient of Variation (%):", coefficient_variation(int_list2))
```

(vi) Write a Python function to find the mean absolute deviation (MAD) of a list of numbers.

Answer:

```
def mad(lst):
    mean_val = np.mean(lst)
    return np.mean([abs(x - mean_val) for x in lst])
print("MAD:", mad(int_list2))
```

(vii) Create a program to calculate the quartile deviation of a list of values.

Answer:

```
def quartile_deviation(lst):
    q1 = np.percentile(lst, 25)
    q3 = np.percentile(lst, 75)
    return (q3 - q1) / 2
print("Quartile Deviation:", quartile_deviation(int_list2))
```

(viii) Implement a function to find the range-based coefficient of dispersion for a dataset.

Answer:

```
def range_dispersion(lst):
    return (max(lst) - min(lst)) / (max(lst) + min(lst))
print("Range-based Coefficient of Dispersion:", range_dispersion(int_list2))
```

Q3. Write a Python class representing a discrete random variable with methods to calculate its expected value and variance.

Answer:

```
class DiscreteRandomVariable:  
    def __init__(self, values, probabilities):  
        self.values = np.array(values)  
        self.probabilities = np.array(probabilities)  
  
    def expected_value(self):  
        return np.sum(self.values * self.probabilities)  
  
    def variance(self):  
        mu = self.expected_value()  
        return np.sum(((self.values - mu) ** 2) * self.probabilities)  
  
drv = DiscreteRandomVariable([1, 2, 3], [0.2, 0.5, 0.3])  
  
print("Expected Value:", drv.expected_value())  
  
print("Variance:", drv.variance())
```

Q4. Implement a program to simulate rolling of a fair six-sided die and calculate expected value and variance.

Answer:

```
np.random.seed(0)  
  
die_rolls = np.random.randint(1, 7, 1000)  
  
print("Expected Value:", np.mean(die_rolls))  
  
print("Variance:", np.var(die_rolls, ddof=1))
```

Q5. Create a Python function to generate random samples from a distribution (Binomial/Poisson) and calculate mean and variance.

Answer:

```
def generate_samples(distribution="binomial", size=1000):  
  
    if distribution == "binomial":  
  
        samples = np.random.binomial(n=10, p=0.5, size=size)  
  
    elif distribution == "poisson":  
  
        samples = np.random.poisson(lam=5, size=size)
```

```
else:  
    return None  
  
return samples, np.mean(samples), np.var(samples, ddof=1)  
  
samples, mean_val, var_val = generate_samples("poisson", 1000)  
  
print("Mean:", mean_val)  
  
print("Variance:", var_val)
```

Q6. Write a script to generate random numbers from Gaussian distribution and compute mean, variance and standard deviation.

Answer:

```
np.random.seed(0)  
  
normal_samples = np.random.normal(loc=0, scale=1, size=1000)  
  
print("Mean:", np.mean(normal_samples))  
  
print("Variance:", np.var(normal_samples, ddof=1))  
  
print("Standard Deviation:", np.std(normal_samples, ddof=1))
```

Q7. Use seaborn library to load tips dataset. Find the following for the columns total_bill and tip:

Q(i). Write a Python function that calculates their skewness.

Answer:

```
import seaborn as sns  
  
from scipy import stats  
  
tips = sns.load_dataset("tips")  
  
def skewness(column):  
  
    return stats.skew(column)  
  
print("Skewness of total_bill:", skewness(tips["total_bill"]))  
  
print("Skewness of tip:", skewness(tips["tip"]))
```

Q(ii). Create a program that determines whether the columns exhibit positive skewness, negative skewness, or approximately symmetric.

Answer:

```
def skew_type(value):
```

```
if value > 0:  
    return "Positive Skewness"  
  
elif value < 0:  
    return "Negative Skewness"  
  
else:  
    return "Approximately Symmetric"
```

```
print("total_bill:", skew_type(skewness(tips["total_bill"])))  
print("tip:", skew_type(skewness(tips["tip"])))
```

Q(iii). Write a function that calculates the covariance between two columns.

Answer:

```
def covariance(col1, col2):  
    return np.cov(col1, col2, ddof=1)[0, 1]  
  
print("Covariance between total_bill and tip:", covariance(tips["total_bill"], tips["tip"]))
```

Q(iv). Implement a Python program that calculates the Pearson correlation coefficient between two columns.

Answer:

```
def pearson_corr(col1, col2):  
    return np.corrcoef(col1, col2)[0, 1]  
  
print("Pearson correlation between total_bill and tip:", pearson_corr(tips["total_bill"], tips["tip"]))
```

Q(v). Write a script to visualize the correlation between two specific columns using scatter plot.

Answer:

```
import matplotlib.pyplot as plt  
  
plt.figure(figsize=(7,5))  
  
plt.scatter(tips["total_bill"], tips["tip"])  
  
plt.title("Scatter Plot: total_bill vs tip")  
  
plt.xlabel("total_bill")
```

```
plt.ylabel("tip")
```

```
plt.show()
```

Q8. Write a Python function to calculate the probability density function (PDF) of a continuous random variable for a given normal distribution.

Answer:

```
import numpy as np
```

```
def normal_pdf(x, mu=0, sigma=1):
```

```
    pdf = (1 / (sigma * np.sqrt(2 * np.pi))) * np.exp(-0.5 * ((x - mu) / sigma) ** 2)
```

```
    return pdf
```

```
print("Normal PDF at x=1:", normal_pdf(1, mu=0, sigma=1))
```

Q9. Create a program to calculate the cumulative distribution function (CDF) of exponential distribution.

Answer:

```
import numpy as np
```

```
def exponential_cdf(x, lam=1):
```

```
    cdf = 1 - np.exp(-lam * x)
```

```
    return cdf
```

```
print("Exponential CDF at x=2:", exponential_cdf(2, lam=1))
```

Q10. Write a Python function to calculate the probability mass function (PMF) of Poisson distribution.

Answer:

```
import numpy as np
```

```
import math
```

```
def poisson_pmf(k, lam=5):
```

```
    pmf = (lam ** k * np.exp(-lam)) / math.factorial(k)
```

```
    return pmf
```

```
print("Poisson PMF at k=3:", poisson_pmf(3, lam=5))
```

Q11. A company wants to test if a new website layout leads to a higher conversion rate (percentage of visitors who make a purchase). They collect data from the old and new layouts to compare?

To generate the data use the following command: `` ` python

```
import numpy as np
```

```
# 50 purchases out of 1000 visitors old_layout = np.array([1] * 50 + [0] * 950)
```

```
# 70 purchases out of 1000 visitors
```

```
new_layout = np.array([1] * 70 + [0] * 930)
```

```
``` Apply z-test to find which layout is successful.
```

**ANSWER:**

```
numpy as np
```

```
from statsmodels.stats.proportion import proportions_ztest
```

```
50 purchases out of 1000 visitors
```

```
old_layout = np.array([1] * 50 + [0] * 950)
```

```
70 purchases out of 1000 import visitors
```

```
new_layout = np.array([1] * 70 + [0] * 930)
```

```
Count successes
```

```
count_old = old_layout.sum()
```

```
count_new = new_layout.sum()
```

```
Total observations
```

```
n_old = len(old_layout)
```

```
n_new = len(new_layout)
```

```
Apply Z-test
```

```
z_stat, p_val = proportions_ztest([count_old, count_new], [n_old, n_new])
```

```
print("Z-statistic:", z_stat)
```

```
print("P-value:", p_val)
```

```
if p_val < 0.05:
```

```
 print("Conclusion: New layout is successful (higher conversion rate).")
```

```
else:
```

```
print("Conclusion: No significant difference between old and new layout.")
```

**Q12. A tutoring service claims that its program improves students' exam scores. A sample of students who participated in the program was taken, and their scores before and after the program were recorded?**

Use the below code to generate samples of respective arrays of marks: ```` python  
before\_program = np.array([75, 80, 85, 70, 90, 78, 92, 88, 82, 87]) after\_program =  
np.array([80, 85, 90, 80, 92, 80, 95, 90, 85, 88]) ```` Use z-test to find if the claims made by tutor are true or false.

**Answer:**

```
import numpy as np
from scipy import stats
before_program = np.array([75, 80, 85, 70, 90, 78, 92, 88, 82, 87])
after_program = np.array([80, 85, 90, 80, 92, 80, 95, 90, 85, 88])
Differences
diff = after_program - before_program
One-sample z-test on differences (assuming population std from sample)
mean_diff = np.mean(diff)
std_diff = np.std(diff, ddof=1)
n = len(diff)
z_stat = mean_diff / (std_diff / np.sqrt(n))
p_val = 1 - stats.norm.cdf(z_stat) # one-tailed test (improvement)
print("Z-statistic:", z_stat)
print("P-value:", p_val)
if p_val < 0.05:
 print("Conclusion: Tutor claim is TRUE (scores improved).")
else:
 print("Conclusion: Tutor claim is FALSE (no significant improvement).")
```

**Q13. A pharmaceutical company wants to determine if a new drug is effective in reducing blood pressure. They conduct a study and record blood pressure measurements before and after administering the drug?**

Use the below code to generate samples of respective arrays of blood pressure:

```
```python before_drug = np.array([145, 150, 140, 135, 155, 160, 152, 148, 130, 138])
after_drug = np.array([130, 140, 132, 128, 145, 148, 138, 136, 125, 130])```
Implement z-test to find if the drug really works or not.
```

ANSWER:

```
import numpy as np
from scipy import stats
before_drug = np.array([145, 150, 140, 135, 155, 160, 152, 148, 130, 138])
after_drug = np.array([130, 140, 132, 128, 145, 148, 138, 136, 125, 130])
# Differences (before - after) because reduction is expected
diff = before_drug - after_drug
# One-sample z-test on differences
mean_diff = np.mean(diff)
std_diff = np.std(diff, ddof=1)
n = len(diff)
z_stat = mean_diff / (std_diff / np.sqrt(n))
p_val = 1 - stats.norm.cdf(z_stat) # one-tailed test (reduction)
print("Z-statistic:", z_stat)
print("P-value:", p_val)
if p_val < 0.05:
    print("Conclusion: Drug works (blood pressure reduced).")
else:
    print("Conclusion: Drug does not work (no significant reduction).")
```

Q14. A customer service department claims that their average response time is less than 5 minutes. A sample of recent customer interactions was taken, and the response times were recorded?

Implement the below code to generate the array of response time: ```python
response_times = np.array([4.3, 3.8, 5.1, 4.9, 4.7, 4.2, 5.2, 4.5, 4.6, 4.4])``` Implement z-test to find the claims made by customer service department are true or false

ANSWER:

```

import numpy as np
from scipy import stats
response_times = np.array([4.3, 3.8, 5.1, 4.9, 4.7, 4.2, 5.2, 4.5, 4.6, 4.4])
# Hypothesis:
# H0: mean >= 5
# H1: mean < 5
sample_mean = np.mean(response_times)
sample_std = np.std(response_times, ddof=1)
n = len(response_times)
# Z-test statistic
z_stat = (sample_mean - 5) / (sample_std / np.sqrt(n))
# One-tailed p-value (left tail)
p_val = stats.norm.cdf(z_stat)
print("Sample Mean:", sample_mean)
print("Z-statistic:", z_stat)
print("P-value:", p_val)
if p_val < 0.05:
    print("Conclusion: Claim is TRUE (average response time < 5 minutes).")
else:
    print("Conclusion: Claim is FALSE (average response time is not less than 5 minutes).")

```

Q15. A company is testing two different website layouts to see which one leads to higher click-through rates. Write a Python function to perform an A/B test analysis, including calculating the t-statistic, degrees of freedom, and p-value?

ANSWER:

```

import numpy as np
from scipy import stats
layout_a_clicks = [28, 32, 33, 29, 31, 34, 30, 35, 36, 37]
layout_b_clicks = [40, 41, 38, 42, 39, 44, 43, 41, 45, 47]

```

```

def ab_test_ttest(a, b):
    a = np.array(a)
    b = np.array(b)

    # Welch's t-test (better when variances may differ)
    t_stat, p_val = stats.ttest_ind(a, b, equal_var=False)

    # Degrees of freedom (Welch-Satterthwaite approximation)
    s1 = np.var(a, ddof=1)
    s2 = np.var(b, ddof=1)
    n1 = len(a)
    n2 = len(b)

    dof = (s1/n1 + s2/n2)**2 / (((s1/n1)**2)/(n1-1) + ((s2/n2)**2)/(n2-1))

    return t_stat, dof, p_val

t_stat, dof, p_val = ab_test_ttest(layout_a_clicks, layout_b_clicks)
print("T-statistic:", t_stat)
print("Degrees of freedom:", dof)
print("P-value:", p_val)

```

Q16. A pharmaceutical company wants to determine if a new drug is more effective than an existing drug in reducing cholesterol levels. Create a program to analyze the clinical trial data and calculate the tstatistic and p-value for the treatment effect?

Use the following data of cholesterol level:

```
python
existing_drug_levels = [180, 182, 175, 185, 178, 176, 172, 184, 179, 183]
new_drug_levels = [170, 172, 165, 168, 175, 173, 170, 178, 172, 176]
```

ANSWER:

```

import numpy as np
from scipy import stats
existing_drug_levels = [180, 182, 175, 185, 178, 176, 172, 184, 179, 183]
new_drug_levels = [170, 172, 165, 168, 175, 173, 170, 178, 172, 176]
# Convert to numpy arrays

```

```

existing = np.array(existing_drug_levels)

new = np.array(new_drug_levels)

# Welch's t-test (does not assume equal variance)

t_stat, p_two_tailed = stats.ttest_ind(new, existing, equal_var=False)

# Convert to one-tailed p-value (since we test new < existing)

p_one_tailed = p_two_tailed / 2 if t_stat < 0 else 1 - (p_two_tailed / 2)

print("t-statistic =", t_stat)

print("p-value (one-tailed) =", p_one_tailed)

```

Q17. A school district introduces an educational intervention program to improve math scores. Write a Python function to analyze pre- and post-intervention test scores, calculating the t-statistic and p-value to determine if the intervention had a significant impact.

Use the following data of test score: ` ` `python pre_intervention_scores = [80, 85, 90, 75, 88, 82, 92, 78, 85, 87] post_intervention_scores = [90, 92, 88, 92, 95, 91, 96, 93, 89, 93]

ANSWER:

```

import numpy as np

from scipy import stats

def intervention_ttest(pre_scores, post_scores):

    pre = np.array(pre_scores)

    post = np.array(post_scores)

    # Paired t-test

    t_stat, p_value = stats.ttest_rel(post, pre)

    return t_stat, p_value

# Given data

pre_intervention_scores = [80, 85, 90, 75, 88, 82, 92, 78, 85, 87]

post_intervention_scores = [90, 92, 88, 92, 95, 91, 96, 93, 89, 93]

t_stat, p_val = intervention_ttest(pre_intervention_scores, post_intervention_scores)

print("t-statistic =", t_stat), print("p-value =", p_val)

```

Q18. An HR department wants to investigate if there's a gender-based salary gap within the company. Develop a program to analyze salary data, calculate the t-statistic, and determine if there's a statistically significant difference between the average salaries of male and female employees.

Use the below code to generate synthetic data:

```
``` python # Generate synthetic salary data for male and female employees
np.random.seed(0) # For reproducibility
male_salaries = np.random.normal(loc=50000, scale=10000, size=20)
female_salaries = np.random.normal(loc=55000, scale=9000, size=20)
```

**ANSWER:**

```
import numpy as np
from scipy import stats

Generate synthetic salary data for male and female employees
np.random.seed(0) # For reproducibility
male_salaries = np.random.normal(loc=50000, scale=10000, size=20)
female_salaries = np.random.normal(loc=55000, scale=9000, size=20)

Welch's t-test (recommended when variances may differ)
t_stat, p_value = stats.ttest_ind(male_salaries, female_salaries, equal_var=False)

print("Male salaries mean :", np.mean(male_salaries))
print("Female salaries mean :", np.mean(female_salaries))

print("t-statistic :", t_stat)
print("p-value :", p_value)

Decision (alpha = 0.05)
alpha = 0.05

if p_value < alpha:
 print("Statistically significant difference in average salaries.")
else:
 print("No statistically significant difference in average salaries.")
```

**Q19. A manufacturer produces two different versions of a product and wants to compare their quality scores. Create a Python function to analyze quality assessment data, calculate the t-statistic, and decide whether there's a significant difference in quality between the two versions.**

Use the following data: `` python version1\_scores = [85, 88, 82, 89, 87, 84, 90, 88, 85, 86, 91, 83, 87, 84, 89, 86, 84, 88, 85, 86, 89, 90, 87, 88, 85] version2\_scores = [80, 78, 83, 81, 79, 82, 76, 80, 78, 81, 77, 82, 80, 79, 82, 79, 80, 81, 79, 82, 79, 78, 80, 81, 82]

**ANSWER:**

```
import numpy as np
from scipy import stats

def compare_quality(version1_scores, version2_scores, alpha=0.05):
 v1 = np.array(version1_scores)
 v2 = np.array(version2_scores)

 # Welch's t-test (safe option: does not assume equal variance)
 t_stat, p_value = stats.ttest_ind(v1, v2, equal_var=False)

 print("Mean Version 1 =", np.mean(v1))
 print("Mean Version 2 =", np.mean(v2))

 print("t-statistic =", t_stat)
 print("p-value =", p_value)

 if p_value < alpha:
 print(" Significant difference in quality between the two versions.")
 else:
 print("No significant difference in quality between the two versions.")

 return t_stat, p_value

Given data
version1_scores = [85, 88, 82, 89, 87, 84, 90, 88, 85, 86, 91, 83, 87, 84, 89, 86, 84, 88, 85, 86, 89, 90, 87, 88, 85]
version2_scores = [80, 78, 83, 81, 79, 82, 76, 80, 78, 81, 77, 82, 80, 79, 82, 79, 80, 81, 79, 82, 79, 80, 81, 82]
compare_quality(version1_scores, version2_scores)
```

**20. A restaurant chain collects customer satisfaction scores for two different branches. Write a program to analyze the scores, calculate the t-statistic, and**

**determine if there's a statistically significant difference in customer satisfaction between the branches.**

Use the below data of scores: ````python branch\_a\_scores = [4, 5, 3, 4, 5, 4, 5, 3, 4, 4, 5, 4, 4, 3, 4, 5, 5, 4, 3, 4, 5, 4, 3, 5, 4, 4, 5, 3, 4, 5, 4] branch\_b\_scores = [3, 4, 2, 3, 4, 3, 4, 2, 3, 3, 4, 3, 3, 2, 3, 4, 4, 3, 2, 3, 4, 3, 2, 4, 3, 3, 4, 2, 3, 4, 3]

**ANSWER:**

```
import numpy as np
from scipy import stats
branch_a_scores = [4, 5, 3, 4, 5, 4, 5, 3, 4, 4, 5, 4, 4, 3, 4, 5, 5, 4, 3, 4, 5, 4, 3, 5, 4, 4, 5, 3, 4, 5, 4]
branch_b_scores = [3, 4, 2, 3, 4, 3, 4, 2, 3, 3, 4, 3, 3, 2, 3, 4, 4, 3, 2, 3, 4, 3, 2, 3, 4, 3, 4, 2, 3, 4, 3]
a = np.array(branch_a_scores)
b = np.array(branch_b_scores)
Welch's t-test (best for real-world data)
t_stat, p_value = stats.ttest_ind(a, b, equal_var=False)
print("Branch A Mean =", np.mean(a))
print("Branch B Mean =", np.mean(b))
print("t-statistic =", t_stat)
print("p-value =", p_value)
Decision
alpha = 0.05
if p_value < alpha:
 print(" Statistically significant difference in satisfaction between branches.")
else:
 print(" No statistically significant difference in satisfaction between branches.")
```

**Q21. A political analyst wants to determine if there is a significant association between age groups and voter preferences (Candidate A or Candidate B). They collect data from a sample of 500 voters and classify them into different age groups and candidate preferences. Perform a Chi-Square test to determine if there is a significant association between age groups and voter preferences?**

Use the below code to generate data:

```
``` python
np.random.seed(0)
age_groups = np.random.choice([18 30 , 31 50 , 51+', 51+'], size=30)
voter_preferences = np.random.choice(['Candidate A', 'Candidate B'], size=30)
```

ANSWER:

```
import numpy as np
import pandas as pd
from scipy.stats import chi2_contingency
np.random.seed(0)

# Generate synthetic data (fixed age group labels)
age_groups = np.random.choice(['18-30', '31-50', '51+'], size=30)
voter_preferences = np.random.choice(['Candidate A', 'Candidate B'], size=30)

# Create contingency table
data = pd.crosstab(age_groups, voter_preferences)
print("Contingency Table:\n", data)

# Chi-square test
chi2, p_value, dof, expected = chi2_contingency(data)
print("\nChi-square statistic =", chi2)
print("Degrees of freedom  =", dof)
print("p-value      =", p_value)

# Decision
alpha = 0.05
if p_value < alpha:
    print(" Significant association between age groups and voter preference.")
else:
    print(" No significant association between age groups and voter preference.")
```

22. A company conducted a customer satisfaction survey to determine if there is a significant relationship between product satisfaction levels (Satisfied, Neutral, Dissatisfied) and the region where customers are located (East, West, North, South). The survey data is summarized in a contingency table. Conduct a

ChiSquare test to determine if there is a significant relationship between product satisfaction levels and customer regions.

Sample data: ````python #Sample data: Product satisfaction levels (rows) vs. Customer regions (columns) data = np.array([[50, 30, 40, 20], [30, 40, 30, 50], [20, 30, 40, 30]])

ANSWER:

```
import numpy as np
from scipy.stats import chi2_contingency
# Sample data: Satisfaction levels (rows) vs Regions (columns)
# Rows: [Satisfied, Neutral, Dissatisfied]
# Cols: [East, West, North, South]
data = np.array([
    [50, 30, 40, 20],
    [30, 40, 30, 50],
    [20, 30, 40, 30]
])
# Chi-square test
chi2, p_value, dof, expected = chi2_contingency(data)
print("Chi-square statistic =", chi2)
print("Degrees of freedom  =", dof)
print("p-value      =", p_value)
print("\nExpected Frequencies:\n", expected)
# Decision
alpha = 0.05
if p_value < alpha:
    print("\n Significant relationship between satisfaction level and region.")
else:
    print("\n No significant relationship between satisfaction level and region.")
```

23. A company implemented an employee training program to improve job performance (Effective, Neutral, Ineffective). After the training, they collected data

from a sample of employees and classified them based on their job performance before and after the training. Perform a Chi-Square test to determine if there is a significant difference between job performance levels before and after the training.

Sample data: ````python # Sample data: Job performance levels before (rows) and after (columns) training data = np.array([[50, 30, 20], [30, 40, 30], [20, 30, 40]])

ANSWER:

```
import numpy as np
from scipy.stats import chi2_contingency
# Sample data: performance levels before (rows) and after (columns)
# Rows: [Effective, Neutral, Ineffective] (Before)
# Columns: [Effective, Neutral, Ineffective] (After)
data = np.array([
    [50, 30, 20],
    [30, 40, 30],
    [20, 30, 40]
])
# Chi-square test
chi2, p_value, dof, expected = chi2_contingency(data)
print("Chi-square statistic =", chi2)
print("Degrees of freedom  =", dof)
print("p-value      =", p_value)
print("\nExpected Frequencies:\n", expected)
# Decision
alpha = 0.05
if p_value < alpha:
    print("\n Significant difference between performance levels before and after training.")
else:
    print("\n No significant difference between performance levels before and after training.")
```

24. A company produces three different versions of a product: Standard, Premium, and Deluxe. The company wants to determine if there is a significant difference in customer satisfaction scores among the three product versions. They conducted a survey and collected customer satisfaction scores for each version from a random sample of customers. Perform an ANOVA test to determine if there is a significant difference in customer satisfaction scores.

Use the following data:

```
```python # Sample data: Customer satisfaction scores for each product version
standard_scores = [80, 85, 90, 78, 88, 82, 92, 78, 85, 87]
premium_scores = [90, 92, 88, 92, 95, 91, 96, 93, 89, 93]
deluxe_scores = [95, 98, 92, 97, 96, 94, 98, 97, 92, 99]
```

**ANSWER:**

```
import numpy as np
from scipy.stats import f_oneway
Sample data
standard_scores = [80, 85, 90, 78, 88, 82, 92, 78, 85, 87]
premium_scores = [90, 92, 88, 92, 95, 91, 96, 93, 89, 93]
deluxe_scores = [95, 98, 92, 97, 96, 94, 98, 97, 92, 99]
One-way ANOVA test
f_stat, p_value = f_oneway(standard_scores, premium_scores, deluxe_scores)
print("F-statistic =", f_stat)
print("p-value ==", p_value)
Decision
alpha = 0.05
if p_value < alpha:
 print(" Significant difference in customer satisfaction among product versions.")
else:
 print("No significant difference in customer satisfaction among product versions.")
```